

## Machine Learning on Kubernetes

### Setting up a functional Kubernetes cluster

1. First, we create a Kubernetes cluster with three nodes if we deleted the existing ones from the previous HomeWorks. So, we enable GKE first and create Kubernetes cluster:

```
gcloud container clusters create kubia --num-nodes=1 --machine-type=e2-micro --region=us-west1
```

```
nseghid8444@cloudshell:~ (cs571-cloud-computing-19967)$ gcloud container clusters create kubia --num-nodes=1 --machine-type=e2-micro --region=us-west1-b
Default change: VPC-native is the default mode during cluster creation for versions greater than 1.21.0-gke.1500. To create advanced routes based clusters, please pass the '--no-enable-ip-alias' flag
Note: Your Pod address range ('--cluster-ipv4-cidr') can accommodate at most 1008 node(s).
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...working..
.
Creating cluster kubia in us-west1-b... Cluster is being health-checked (master is healthy)...done.
Created [https://container.googleapis.com/v1/projects/cs571-cloud-computing-19967/zones/us-west1-b/clusters/kubia].
To inspect the contents of your cluster, go to: https://console.cloud.google.com/kubernetes/workload/_gcloud/us-west1-b/kubia?project=cs571-cloud-computing-19967
kubeconfig entry generated for kubia.
NAME: kubia
LOCATION: us-west1-b
MASTER_VERSION: 1.27.8-gke.1067004
MASTER_IP: 34.127.47.41
MACHINE_TYPE: e2-micro
NODE_VERSION: 1.27.8-gke.1067004
NUM_NODES: 1
STATUS: RUNNING
nseghid8444@cloudshell:~ (cs571-cloud-computing-19967)$
```

2. Check if the cluster and the nodes are correctly created and running.
  - Make sure your GKE cluster is running by: `gcloud container clusters list`

```
nseghid8444@cloudshell:~ (cs571-cloud-computing-19967)$ gcloud container clusters list
NAME: kubia
LOCATION: us-west1-b
MASTER_VERSION: 1.27.8-gke.1067004
MASTER_IP: 34.127.47.41
MACHINE_TYPE: e2-micro
NODE_VERSION: 1.27.8-gke.1067004
NUM_NODES: 1
STATUS: RUNNING
nseghid8444@cloudshell:~ (cs571-cloud-computing-19967)$
```

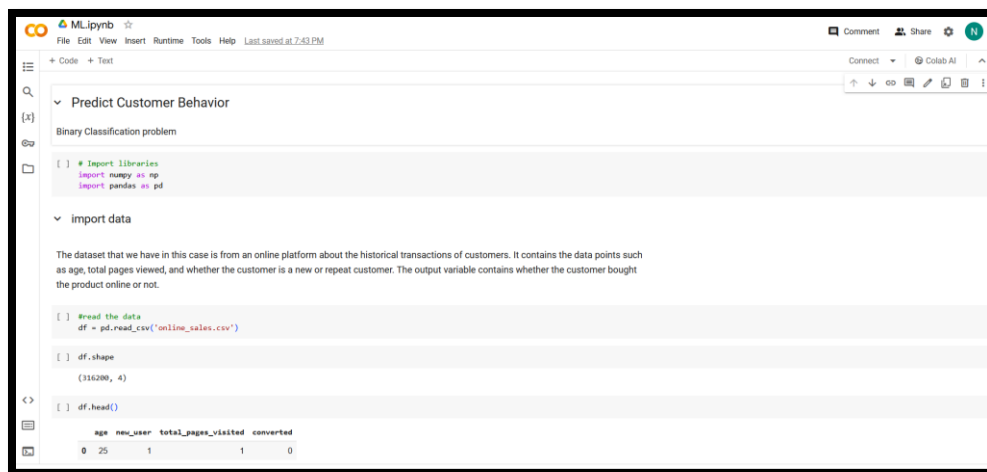
- Or we can use minikube to run a single node using: `minikube start`

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ minikube start
* minikube v1.32.0 on Debian 11.9 (amd64)
- MINIKUBE_FORCE_SYSTEMD=true
- MINIKUBE_HOME=/google/minikube
- MINIKUBE_WANTUPDATENOTIFICATION=false
* Using the docker driver based on existing profile
* Starting control plane node minikube in cluster minikube
* Pulling base image ...
* Updating the running docker "minikube" container ...
```

## Implementing the procedures described on Chapter 4 Machine Learning Deployment Using Docker

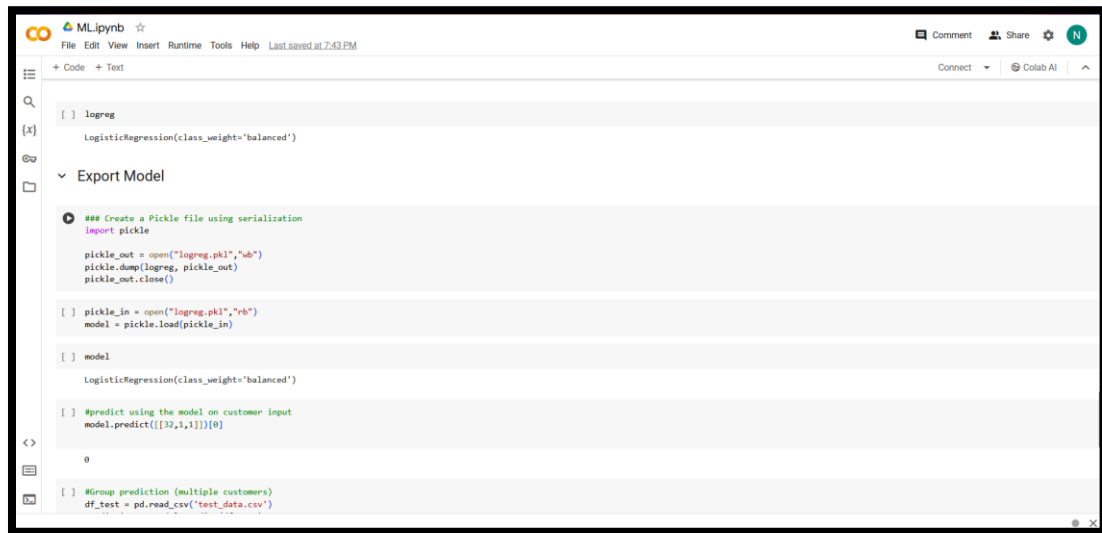
### Step 1. Train a machine learning model

- Build, train, and test a machine learning model.



### Step 2: Save and export the trained ML model.

- Once The model is trained, we need to save it using pickle/joblib in order to reuse it later during predictions.



### Step 3: Create a Flask app including the UI layer

1. Here we will build a Flask app along with Flasgger to deploy an ML model easily. We will do this by creating a flask\_api.py file with the following content.

```
# -*- coding: utf-8 -*-  
"""  
Created on Mon May 25 12:50:04 2020  
@author: pramod.singh  
"""  
from flask import Flask, request  
import numpy as np  
import pickle  
import pandas as pd  
from flasgger import Swagger  
  
app = Flask(__name__)  
Swagger(app)  
  
pickle_in = open("logreg.pkl", "rb")  
model = pickle.load(pickle_in)  
  
@app.route('/')  
def home():  
    return "Welcome to the Flask API!"  
  
@app.route('/predict', methods=["GET"])  
def predict_class():  
    """Predict if Customer would buy the product or not.  
    ---
```

## Niyat Habtom

```
parameters:
  - name: age
    in: query
    type: number
    required: true
  - name: new_user
    in: query
    type: number
    required: true
  - name: total_pages_visited
    in: query
    type: number
    required: true
responses:
  200:
    description: Prediction
    """
    age = int(request.args.get("age"))
    new_user = int(request.args.get("new_user"))
    total_pages_visited = int(request.args.get("total_pages_visited"))
    prediction = model.predict([[age, new_user, total_pages_visited]])
    return "Model prediction is " + str(prediction)

@app.route('/predict_file', methods=["POST"])
def prediction_test_file():
    """Prediction on multiple input test file.
    ---
    parameters:
      - name: file
        in: formData
        type: file
        required: true
    responses:
      200:
        description: Test file Prediction
        """
    df_test = pd.read_csv(request.files.get("file"))
    prediction = model.predict(df_test)
    return str(list(prediction))

if __name__ == '__main__':
    app.run(debug=True, host='0.0.0.0', port=5000)
```

## Niyat Habtom

```
flask_api.py X
C: > Users > Niyat Habtom Seghid > Downloads > flask_api.py
1  # -*- coding: utf-8 -*-
2  """
3  Created on Mon May 25 12:50:04 2020
4
5  @author: pramod.singh
6  """
7
8  from flask import Flask, request
9  import numpy as np
10 import pickle
11 import pandas as pd
12 import flasgger
13 from flasgger import Swagger
14
15 app=Flask(__name__)
16 Swagger(app)
17
18 pickle_in = open("logreg.pkl", "rb")
19 model=pickle.load(pickle_in)
20
21
22 @app.route('/predict', methods=["Get"])
23 def predict_class():
24
25     """Predict if Customer would buy the product or not .
26     ---
27     parameters:
```

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ cat flask_api.py
# -*- coding: utf-8 -*-
"""
Created on Mon May 25 12:50:04 2020
@author: pramod.singh
"""
from flask import Flask, request
import numpy as np
import pickle
import pandas as pd
from flasgger import Swagger

app = Flask(__name__)
Swagger(app)

pickle_in = open("logreg.pkl", "rb")
model = pickle.load(pickle_in)

@app.route('/')
def home():
    return "Welcome to the Flask API!"

@app.route('/predict', methods=["GET"])
def predict_class():
    """Predict if Customer would buy the product or not.
    ---
    parameters:
      - name: age
        in: query
        type: number
        required: true
      - name: new_user
        in: query
        type: number
        required: true
      - name: total_pages_visited
        in: query
```

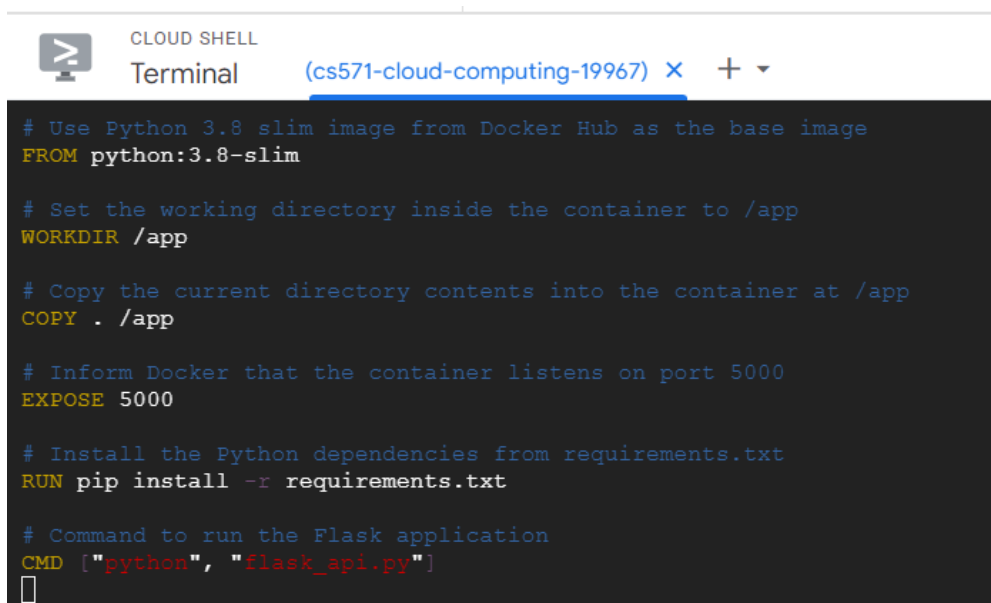
#### Step 4: Creating a custom Docker File for the app

In this section, we will build a classification model from scratch and try to deploy it using a Flask app. The only difference from the previous chapter will be to Dockerize the entire application and run it on any platform. We will also make use of the library Flasgger to handle the UI part of the app to make it more intuitive to consume the results.

1. To start with, let's create a new folder in the local system called **Docker**.

```
nseghid8444@cloudshell:~ (cs571-cloud-computing-19967) $ mkdir Docker
nseghid8444@cloudshell:~ (cs571-cloud-computing-19967) $ cd Docker
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967) $ sudo vim Dockerfile
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967) $
```

2. We are now going to create a **Dockerfile** and mention all the steps to run this app using Docker. This Dockerfile creates a Docker image for a Flask web application: Edit Dockerfile to this new content using command: `sudo vim Dockerfile`



The screenshot shows a terminal window titled 'CLOUD SHELL Terminal' with a tab for '(cs571-cloud-computing-19967)'. The terminal displays the following Dockerfile content:

```
# Use Python 3.8 slim image from Docker Hub as the base image
FROM python:3.8-slim

# Set the working directory inside the container to /app
WORKDIR /app

# Copy the current directory contents into the container at /app
COPY . /app

# Inform Docker that the container listens on port 5000
EXPOSE 5000

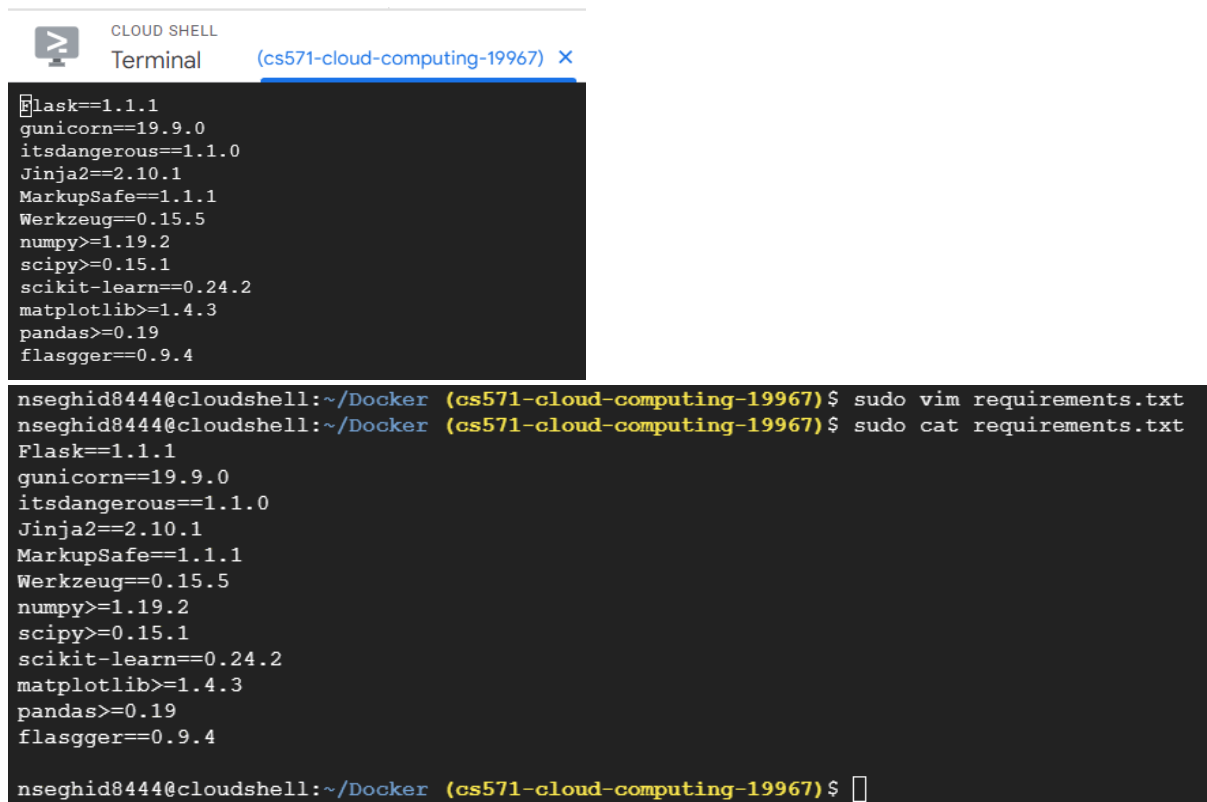
# Install the Python dependencies from requirements.txt
RUN pip install -r requirements.txt

# Command to run the Flask application
CMD ["python", "flask_api.py"]
```

- We start with providing the base image first to the Docker server that needs to be pulled from Docker Hub.
- The next command is to change Docker's current working directory to the directory where we will copy all the files.
- In the next step, we copy all the files and content from the local directory to the Docker directory.
- The next command is to expose port 5000 of Docker to run this application.

- The next step is to install all the dependencies and required libraries.
- The last command in the Dockerfile is the startup command.

3. Create the **requirements.txt** file with the following contents. This file contains all the dependencies and libraries to be installed.



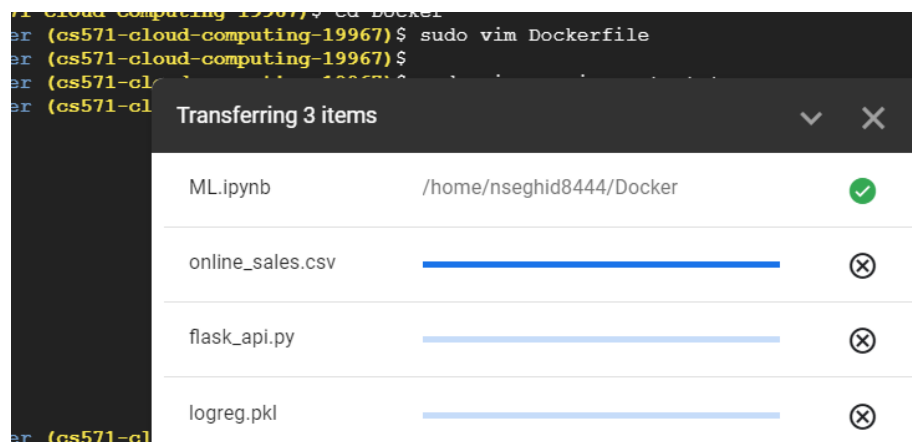
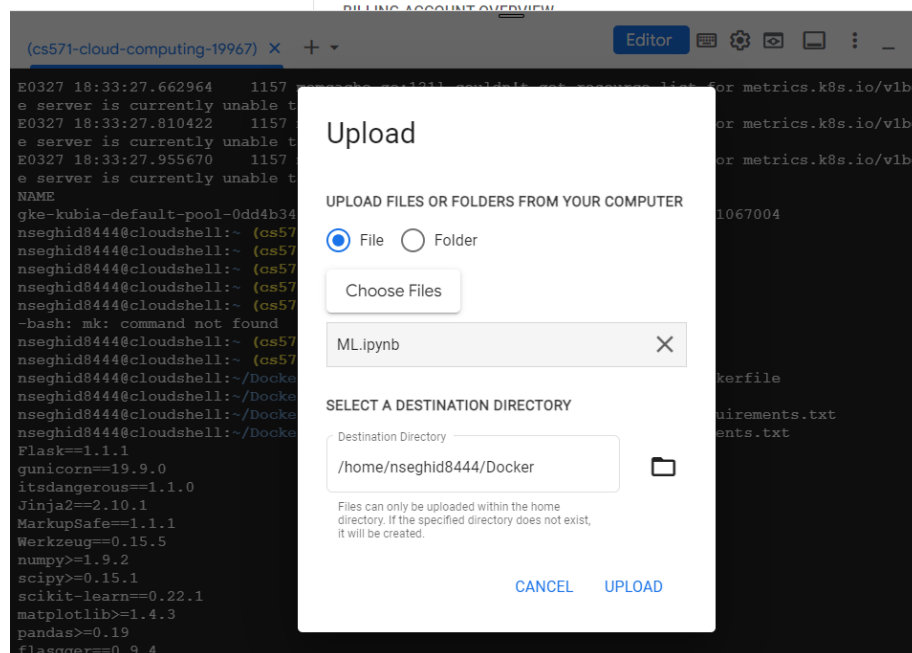
```
CLOUD SHELL
Terminal (cs571-cloud-computing-19967) X

Flask==1.1.1
gunicorn==19.9.0
itsdangerous==1.1.0
Jinja2==2.10.1
MarkupSafe==1.1.1
Werkzeug==0.15.5
numpy>=1.19.2
scipy>=0.15.1
scikit-learn==0.24.2
matplotlib>=1.4.3
pandas>=0.19
flasgger==0.9.4

nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ sudo vim requirements.txt
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ sudo cat requirements.txt
Flask==1.1.1
gunicorn==19.9.0
itsdangerous==1.1.0
Jinja2==2.10.1
MarkupSafe==1.1.1
Werkzeug==0.15.5
numpy>=1.19.2
scipy>=0.15.1
scikit-learn==0.24.2
matplotlib>=1.4.3
pandas>=0.19
flasgger==0.9.4

nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$
```

4. Then we upload all the files we created before to the terminal. Upload logreg.pkl, ML.ipynb, and flask\_api.py files by clicking the three dots at the top-right of the Cloud Shell Terminal and then choose upload.



```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ ls
Dockerfile flask_api.py logreg.pkl ML.ipynb online_sales.csv requirements.txt
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$
```

## Step 5. Run the app using a Docker container

In this step of the process, we build the Docker custom image from the Dockerfile in the previous step and run the container. Once all the commands in the Dockerfile get executed and we have the final image built, we can initiate the container to run our ML app.



1. We have to go to the terminal in the same directory where all the files are present and run the docker build command to build the Docker image from Dockerfile.

```
sudo docker build -t ml_app_docker .
```

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ docker build -t ml_app_docker .
[+] Building 43.2s (9/9) FINISHED
=> [internal] load build definition from Dockerfile                                docker:default 0.0s
=> => transferring dockerfile: 163B                                              0.0s
=> [internal] load metadata for docker.io/library/python:3.8-slim                0.6s
=> [internal] load .dockerignore                                                 0.0s
=> => transferring context: 2B                                                    0.0s
=> [1/4] FROM docker.io/library/python:3.8-slim@sha256:72ae14e80c21f274f31111debd505d8f 3.8s
=> => resolve docker.io/library/python:3.8-slim@sha256:72ae14e80c21f274f31111debd505d8f 0.0s
=> => sha256:72ae14e80c21f274f31111debd505d8fa64536fdf41b57f03930b3baf8 1.86kB / 1.86kB 0.0s
=> => sha256:0ad295b2b84581b1348fa9cad80ea49c4159469e8af9749159d2151ea2 1.37kB / 1.37kB 0.0s
=> => sha256:04977f08feb15b05b809d6547d2ecc9e74e082ac9f9b9b6e6ae2ab9075 6.97kB / 6.97kB 0.0s
=> => sha256:8a1e25ce7c4f75e372e9884f8f7b1bedcfe4a7a7d452eb4b0a1c7477 29.12MB / 29.12MB 0.4s
=> => sha256:1103112ebfc46e01c0f35f3586e5a39c6a9ffa32c1a362d4d5f20e3783 3.51MB / 3.51MB 0.2s
=> => sha256:93d3f6d14ae5338f6f639a4ed5946980d38c016a537a330f20921c5c 11.67MB / 11.67MB 0.3s
=> => sha256:46996c1c5ef3592977cd1c8454cf833bf486a5be36f71847794d97bac47a35 246B / 246B 0.4s
=> => sha256:18dacb59e6d34eadfba0da78f1b3a5f5addfccd45ee854f6af9877b9ed 3.13MB / 3.13MB 0.5s
=> => extracting sha256:8a1e25ce7c4f75e372e9884f8f7b1bedcfe4a7a7d452eb4b0a1c7477c9a9034 1.8s
=> => extracting sha256:1103112ebfc46e01c0f35f3586e5a39c6a9ffa32c1a362d4d5f20e3783c6fdd 0.2s
=> => extracting sha256:93d3f6d14ae5338f6f639a4ed5946980d38c016a537a330f20921c5c7e3995a 0.6s
=> => extracting sha256:46996c1c5ef3592977cd1c8454cf833bf486a5be36f71847794d97bac47a35f 0.0s
=> => extracting sha256:18dacb59e6d34eadfba0da78f1b3a5f5addfccd45ee854f6af9877b9ed5c4b3 0.3s
=> [internal] load build context                                                  0.1s
=> => transferring context: 2.91MB                                              0.1s
=> [2/4] WORKDIR /app                                                            0.5s
=> [3/4] COPY . /app                                                            0.0s
=> [4/4] RUN pip install -r requirements.txt                                    34.1s
=> exporting to image                                                            3.9s
=> => exporting layers                                                            3.9s
=> => writing image sha256:d60ebabd96a3b5bf091a967d186fac15dc069af442803addc990248c52c7 0.0s
=> => naming to docker.io/library/ml_app_docker                                0.0s
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$
```

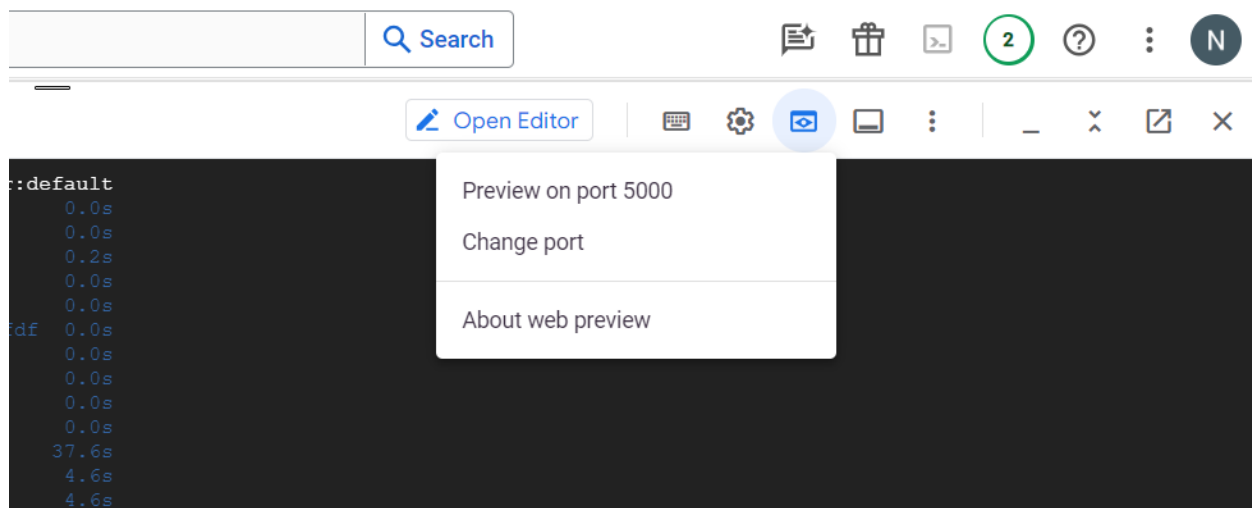
2. Once all the commands in the Dockerfile get executed and we have the final image built from the Dockerfile, we can initiate the container to run our ML app. This command runs a Docker container from the ml\_app\_docker image:

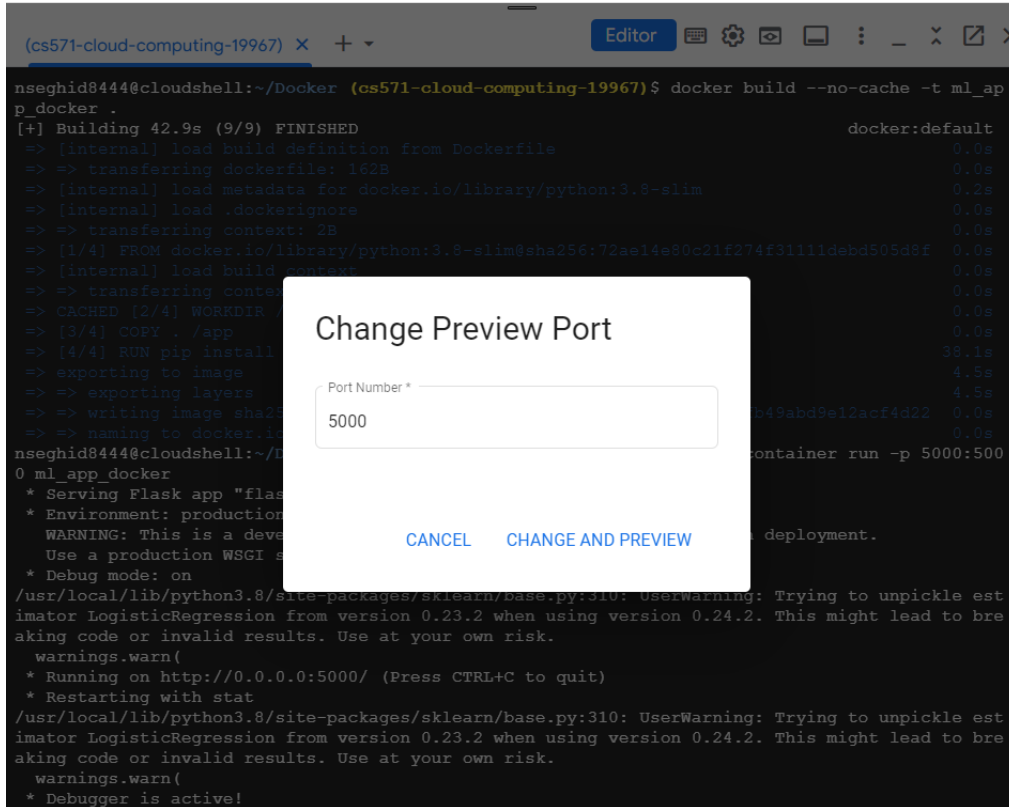
```
sudo docker container run -p 5000:5000 ml_app_docker
```

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ docker container run -p 5000:5000 ml_app_docker
* Serving Flask app "flask_api" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: on
/usr/local/lib/python3.8/site-packages/sklearn/base.py:310: UserWarning: Trying to unpickle estimator LogisticRegression from version 0.23.2 when using version 0.24.2. This might lead to breaking code or invalid results. Use at your own risk.
  warnings.warn(
* Running on http://0.0.0.0:5000/ (Press CTRL+C to quit)
* Restarting with stat
/usr/local/lib/python3.8/site-packages/sklearn/base.py:310: UserWarning: Trying to unpickle estimator LogisticRegression from version 0.23.2 when using version 0.24.2. This might lead to breaking code or invalid results. Use at your own risk.
  warnings.warn(
* Debugger is active!
* Debugger PIN: 778-797-757
```

- We can see that the app has started up successfully, and everything is running inside a Docker container.

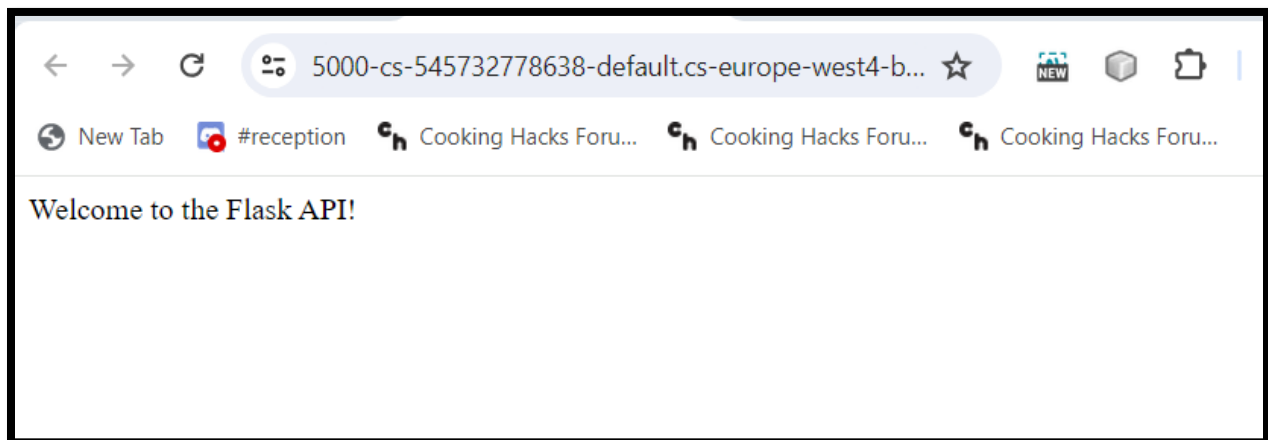
3. In the right-upper side of the terminal, click the web preview option from the menu and then click Preview on port 5000. Change port if it is not 5000 by default.



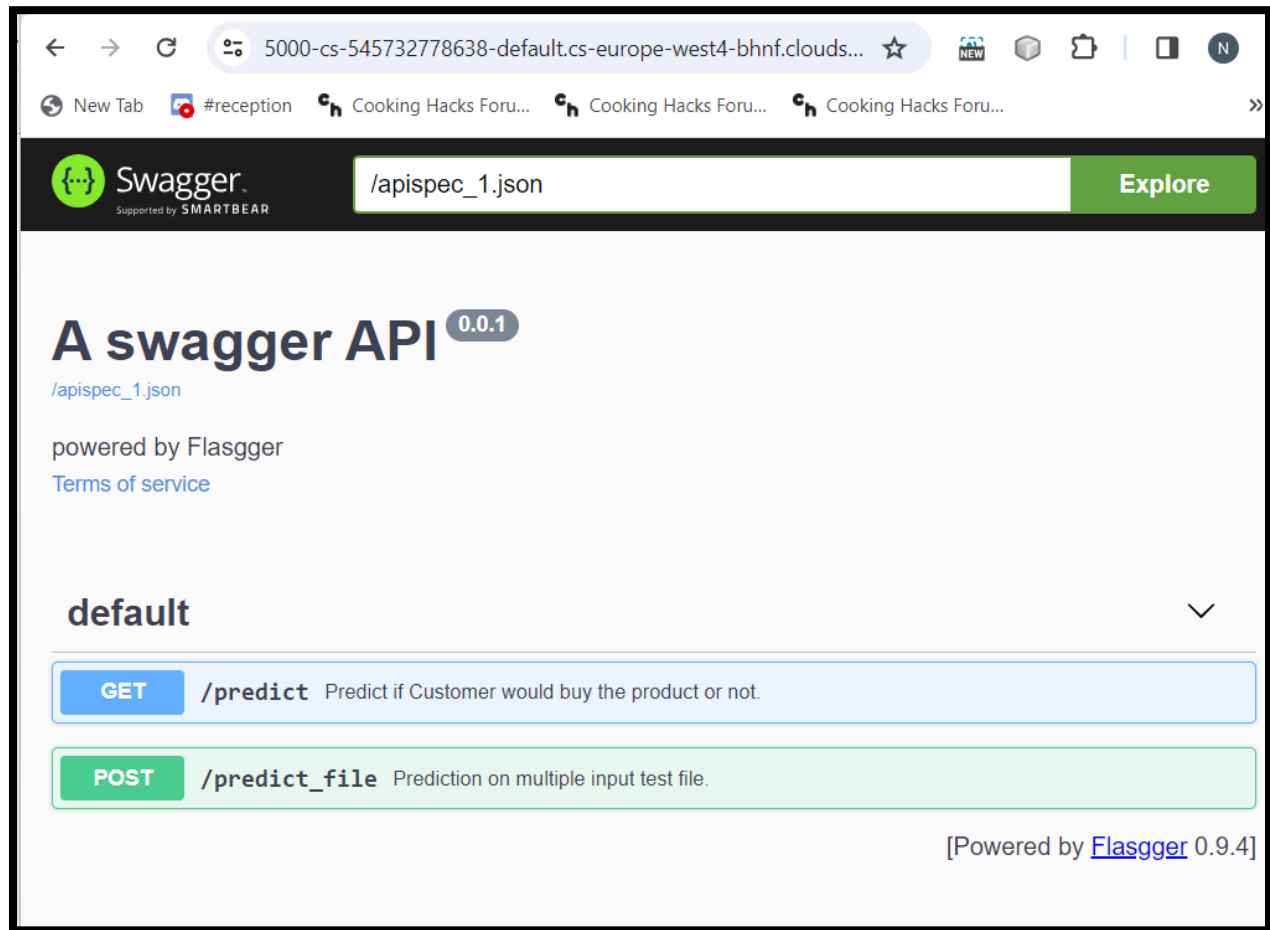


The screenshot shows a terminal window with a dark background. At the top, there's a tab labeled '(cs571-cloud-computing-19967)' and a button labeled 'Editor'. The terminal text shows a Docker build process for a container named 'ml\_app\_docker'. The build steps include loading the Dockerfile, transferring it, loading metadata for 'docker.io/library/python:3.8-slim', and then building the image. A dialog box titled 'Change Preview Port' is overlaid on the terminal. It has a text input field labeled 'Port Number \*' with the value '5000' entered. Below the input field are two buttons: 'CANCEL' and 'CHANGE AND PREVIEW'. The terminal text continues with 'container run -p 5000:5000' and shows the container starting a Flask app. It also displays several warnings from sklearn regarding unpickling estimators.

4. After clicking on Change and Preview you will see this message in the web preview.



5. Add /apidocs/ at the end of the link to access to load the Swagger UI page.



There are two tabs GET and POST. The prediction based on get requests is applicable for single-customer predictions, whereas the Post tab is for the test data prediction (customer group).

- Once we click the Get tab, we can see the options to provide input parameters on which the prediction needs to be made.
- The top-right corner contains a “Try it out” tab that allows us to fill in the values for the input parameters.

6. Click GET and then click Try it out in the top-right corner of the GET box.

The image shows the Swagger UI for a REST API. The browser address bar displays the URL: 5000-cs-545732778638-default.cs-europe-west4-bhnf.clouds... The Swagger logo is in the top left, and the API name is /apispec\_1.json. A green 'Explore' button is in the top right. The main heading is 'A swagger API' with version '0.0.1'. Below it, it says 'powered by Flasgger' and 'Terms of service'. The 'default' section is expanded, showing a 'GET /predict' endpoint with the description 'Predict if Customer would buy the product or not.' The 'Parameters' section is active, showing three query parameters: 'age' (number, required), 'new\_user' (number, required), and 'total\_pages\_visited' (number, required). Each parameter has a corresponding input field. A 'Try it out' button is in the top right of the parameters section.

Swagger  
Supported by SMARTBEAR

/apispec\_1.json

Explore

## A swagger API 0.0.1

/apispec\_1.json

powered by Flasgger  
[Terms of service](#)

### default

**GET** /predict Predict if Customer would buy the product or not.

Parameters Try it out

Name	Description
<b>age</b> * required number (query)	age
<b>new_user</b> * required number (query)	new_user
<b>total_pages_visited</b> * required number	total_pages_visited

This image shows the same Swagger UI interface as the previous one, but with the 'Execute' button highlighted in blue. The 'Parameters' section is still active, showing the three query parameters: 'age', 'new\_user', and 'total\_pages\_visited'. A 'Cancel' button is in the top right of the parameters section. Below the parameters, the 'Execute' button is prominent. The 'Responses' section is also visible, showing a '200' status code with the description 'Prediction'. The 'Response content type' is set to 'application/json'.

**GET** /predict Predict if Customer would buy the product or not.

Parameters Cancel

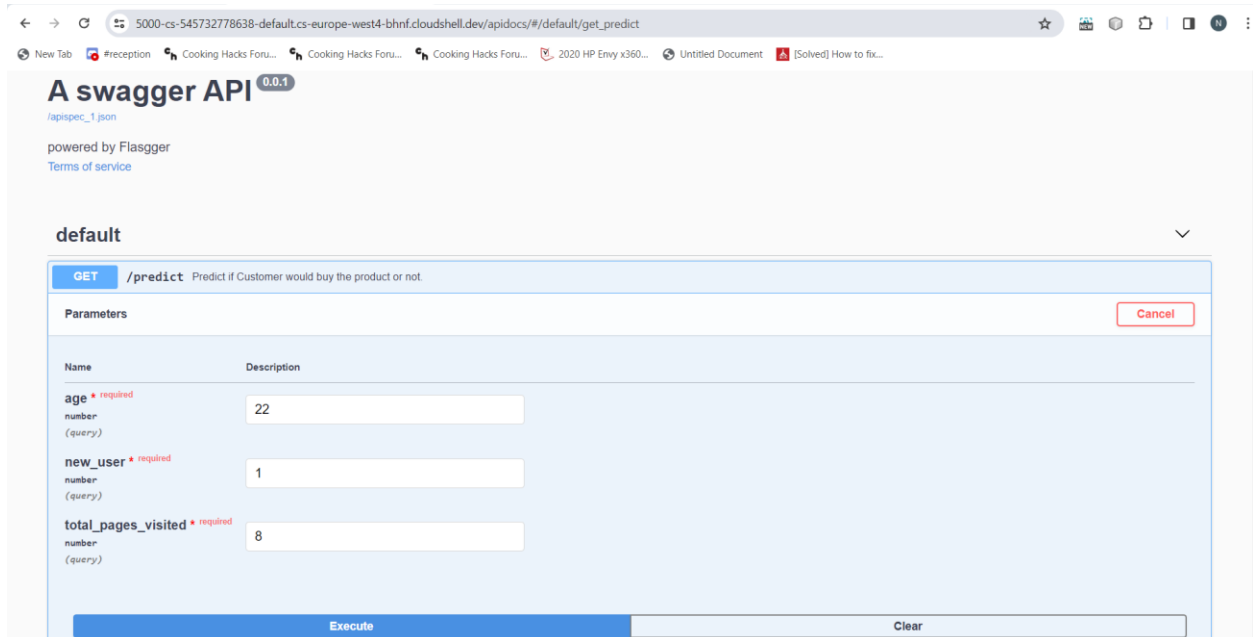
Name	Description
<b>age</b> * required number (query)	age
<b>new_user</b> * required number (query)	new_user
<b>total_pages_visited</b> * required number (query)	total_pages_visited

**Execute**

Responses Response content type: application/json

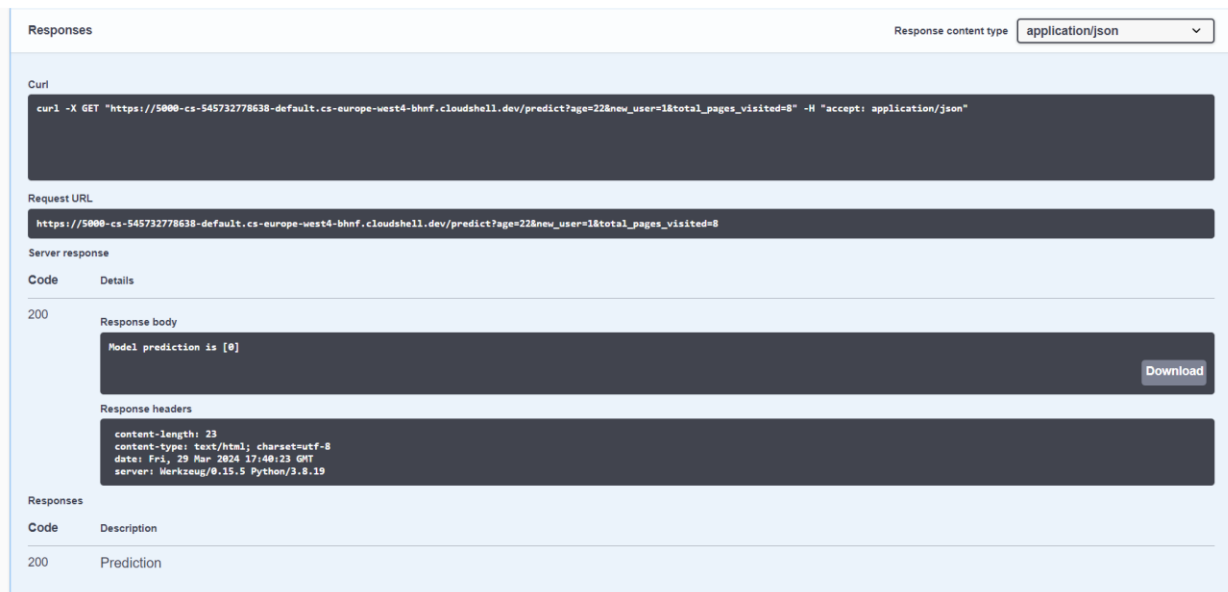
Code	Description
200	Prediction

7. Fill values for the input parameters and then click Execute.



The image shows a Swagger API interface for a service named 'default'. The endpoint is 'GET /predict' with the description 'Predict if Customer would buy the product or not'. The interface includes a 'Parameters' section with three input fields: 'age' (required, number, query) with value '22', 'new\_user' (required, number, query) with value '1', and 'total\_pages\_visited' (required, number, query) with value '8'. There are 'Execute' and 'Clear' buttons at the bottom of the parameters section. The interface is powered by Flasgger and has a version of 0.0.1.

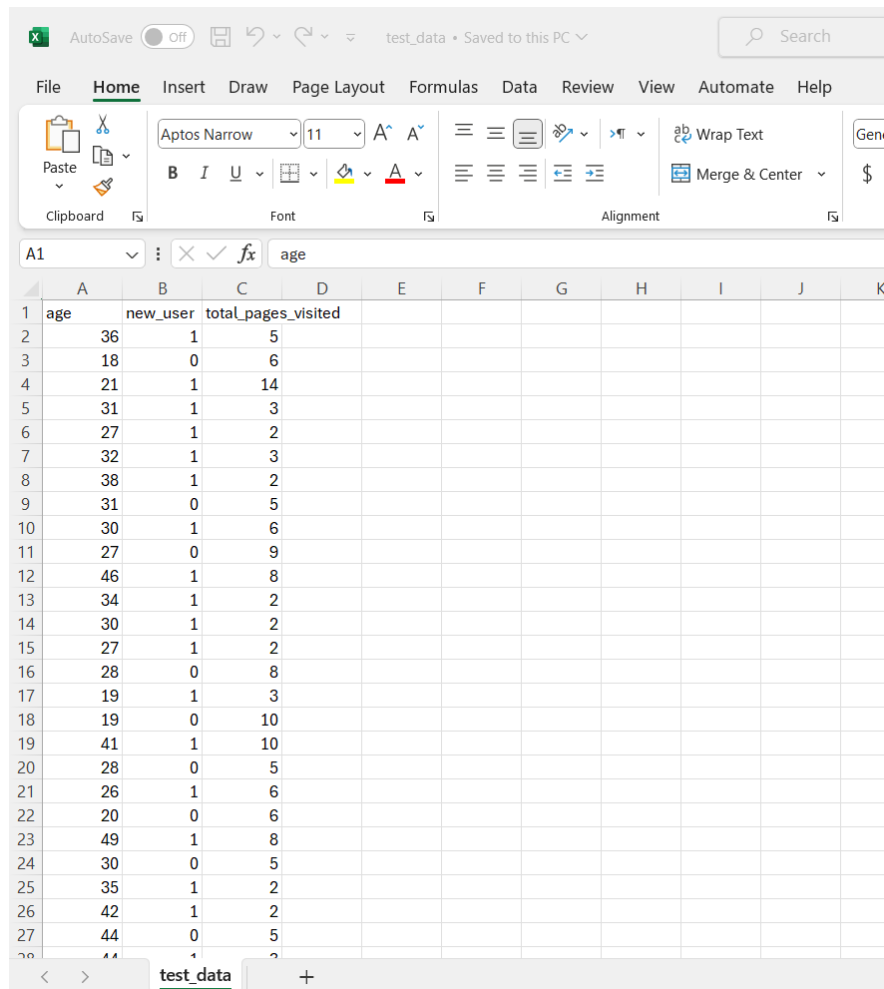
8. Upon the execution call, the request goes to the app, and predictions are made by the model. The result of the model prediction is displayed in the Prediction section of the page as follows.



The image shows the 'Responses' section of the Swagger API interface. The 'Response content type' is set to 'application/json'. The 'Curl' section shows the command: `curl -X GET "https://5000-cs-545732778638-default.cs-europe-west4-bhnf.cloudshell.dev/predict?age=22&new_user=1&total_pages_visited=8" -H "accept: application/json"`. The 'Request URL' is `https://5000-cs-545732778638-default.cs-europe-west4-bhnf.cloudshell.dev/predict?age=22&new_user=1&total_pages_visited=8`. The 'Server response' section shows a '200' status code. The 'Response body' is `Model prediction is [0]`. The 'Response headers' are: `content-length: 23`, `content-type: text/html; charset=utf-8`, `date: Fri, 29 Mar 2024 17:48:23 GMT`, and `server: Werkzeug/0.15.5 Python/3.8.19`. The 'Responses' table at the bottom shows a '200' status code and a 'Prediction' description.

9. Next we have a test\_data csv file that we will use to make predictions for a group of customers through a post request.

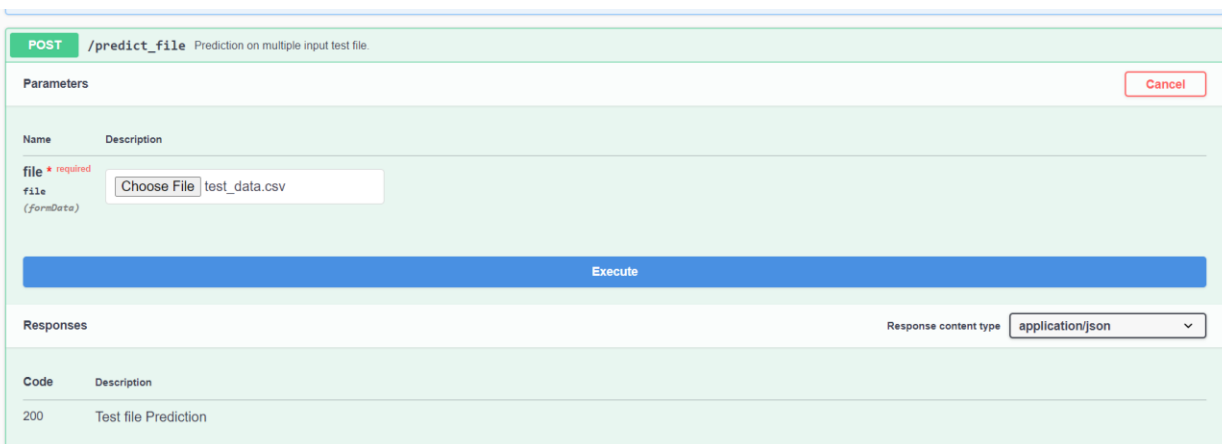
## Niyat Habtom



The screenshot shows a Microsoft Excel spreadsheet with the following data:

	A	B	C	D	E	F	G	H	I	J	K
1	age	new_user	total_pages_visited								
2	36	1	5								
3	18	0	6								
4	21	1	14								
5	31	1	3								
6	27	1	2								
7	32	1	3								
8	38	1	2								
9	31	0	5								
10	30	1	6								
11	27	0	9								
12	46	1	8								
13	34	1	2								
14	30	1	2								
15	27	1	2								
16	28	0	8								
17	19	1	3								
18	19	0	10								
19	41	1	10								
20	28	0	5								
21	26	1	6								
22	20	0	6								
23	49	1	8								
24	30	0	5								
25	35	1	2								
26	42	1	2								
27	44	0	5								
28	44	1	2								

10. Upload the test data using the following order. Click → Post→Try it Out→Choose File.



The screenshot shows a REST client interface with the following details:

- Method:** POST
- URL:** /predict\_file
- Description:** Prediction on multiple input test file.
- Parameters:**
  - Name:** file (required)
  - Description:** file (formData)
  - Value:** Choose File test\_data.csv
- Execute:** A large blue button to execute the request.
- Responses:**
  - Response content type:** application/json
- Code:** 200
- Description:** Test file Prediction

11. Then click Execute and the model would make the predictions. The results would be displayed upon execution as follows.

Name	Description
<b>file</b> * required file <small>(formData)</small>	<input type="text" value="Choose File"/> test_data.csv
<div>ExecuteClear</div>	
ResponsesResponse content typeapplication/json▼	
Curl	<pre>curl -X POST "https://5000-cs-545732778638-default.cs-europe-west4-bnbf.cloudshell.dev/predict_file" -H "accept: application/json" -H "Content-Type: multipart/form-data" -F "file=@test_data.csv;type=csv"</pre>
Request URL	<pre>https://5000-cs-545732778638-default.cs-europe-west4-bnbf.cloudshell.dev/predict_file</pre>
Server response	
CodeDetails	
200	<div>Response body<pre>[{"0": 0, "1": 0, "2": 0, "3": 0, "4": 0, "5": 0, "6": 0, "7": 0, "8": 0, "9": 0, "10": 0, "11": 0, "12": 0, "13": 0, "14": 0, "15": 0, "16": 0, "17": 0, "18": 0, "19": 0, "20": 0, "21": 0, "22": 0, "23": 0, "24": 0, "25": 0, "26": 0, "27": 0, "28": 0, "29": 0, "30": 0, "31": 0, "32": 0, "33": 0, "34": 0, "35": 0, "36": 0, "37": 0, "38": 0, "39": 0, "40": 0, "41": 0, "42": 0, "43": 0, "44": 0, "45": 0, "46": 0, "47": 0, "48": 0, "49": 0, "50": 0, "51": 0, "52": 0, "53": 0, "54": 0, "55": 0, "56": 0, "57": 0, "58": 0, "59": 0, "60": 0, "61": 0, "62": 0, "63": 0, "64": 0, "65": 0, "66": 0, "67": 0, "68": 0, "69": 0, "70": 0, "71": 0, "72": 0, "73": 0, "74": 0, "75": 0, "76": 0, "77": 0, "78": 0, "79": 0, "80": 0, "81": 0, "82": 0, "83": 0, "84": 0, "85": 0, "86": 0, "87": 0, "88": 0, "89": 0, "90": 0, "91": 0, "92": 0, "93": 0, "94": 0, "95": 0, "96": 0, "97": 0, "98": 0, "99": 0}]</pre></div> <div>Download</div>
Response headers	<pre>access-control-allow-credentials: true access-control-allow-methods: GET,POST,OPTIONS,PATCH,DELETE access-control-allow-origin: https://5000-cs-545732778638-default.cs-europe-west4-bnbf.cloudshell.dev content-length: 550 content-type: text/html; charset=utf-8 date: Fri, 29 Mar 2024 12:48:45 GMT server: Werkzeug/0.15.5 Python/3.8.19</pre>
Responses	
CodeDescription	

## Step 6: Stopping/killing the running container.

The last step left after running the application is to stop the running container. This can be done using the `docker stop` or `kill` command on the running container.

1. To see the list of running containers, use the `docker ps` command and select the running container ID to stop it.

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
d9e6d1ca9587	gcr.io/k8s-minikube/kicbase:v0.0.42	"usr/local/bin/entr..."	47 minutes ago	Up 47 minutes	127.0.0.1:
.0.0.1:32769->8443/tcp, 127.0.0.1:32768->32443/tcp	minikube				

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$
```

2. Use the command - `docker kill` to kill the running container as follows.

```
docker kill <Container ID>
```

```
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ docker kill d9e6d1ca9587
d9e6d1ca9587
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$ docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS          NAMES
nseghid8444@cloudshell:~/Docker (cs571-cloud-computing-19967)$
```

## Updated Portfolio- Link to GitHub

<https://github.com/niyat33/Cloud-Computing/blob/main/Kubernetes/Machine%20Learning/README.md>